

Model diagnostics

ML workflow

Problem scoping

Experimentation

- Choose architecture (data, model)
- Train model
- Evaluate model

Deployment

Debugging your ML model


$$J = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)}) + \lambda \sum_{j=1}^n w_j^2$$

Got large prediction errors. What do you do?

- Get more training examples
- Try smaller sets of features
- Try getting additional features
- Try adding polynomial features (e.g., x_1^2 , $x_1 x_2$, x_2^2)
- Try decreasing λ
- Try increasing λ

Model evaluation/selection

Evaluating your model: plotting



Wiggly shape ➡ overfit ➡ poor generalization

Plotting for evaluation is not scalable for high-dimensional data

Evaluating your model: Training/test sets

Split the data into *two* sets:

- **Training set**: the data on which the model is trained
- **Test set**: the data reserved for evaluation. **New to the model**

e.g., 80/20 split: 80% training, 20% test

Computing training and test errors: linear regression example

Fit the model parameters by minimizing the cost function:

$$J = \underbrace{\frac{1}{m_{train}} \sum_{i=1}^{m_{train}} (y - f(x))^2}_{\text{MSE}_{train}} + \lambda \sum_{j=1}^n w_j^2$$

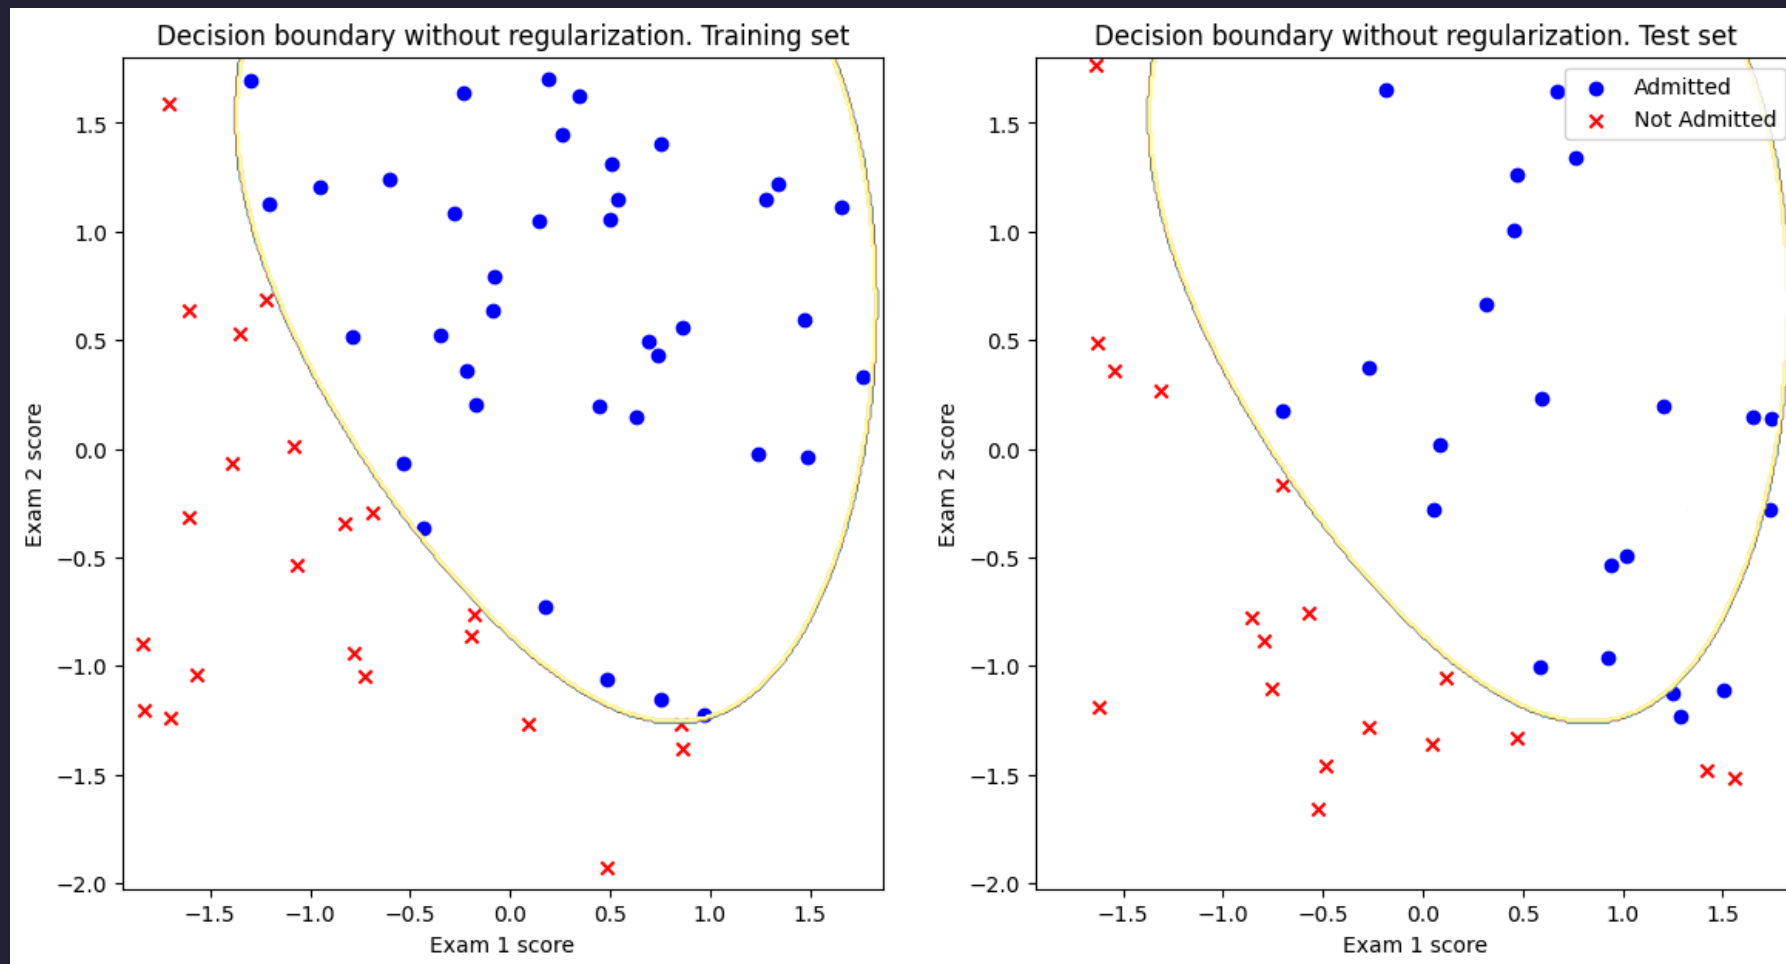
Compute training error:

$$J_{train} = \text{MSE}_{train}$$

Compute test error:

$$J_{test} = \text{MSE}_{test}$$

Overfitting: exam dataset example



J_{train} is low, J_{test} is high

Model selection: Choosing model based on test error

d : degree of polynomial

$$d = 1, f(x) = w_1x + b$$

$$d = 2, f(x) = w_1x + w_2x^2 + b$$

$$d = 3, f(x) = w_1x + w_2x^2 + w_3x^3 + b$$

...

$$d = 10, f(x) = w_1x + w_2x^2 + \dots + w_{10}x^{10} + b$$

Calculate J_{test} for each model (i.e., each d)

Choose the model with the lowest J_{test}


J_{test} is likely to be an optimistic estimate of generalization error because an extra parameter d was chosen using the test set

Training/validation/test sets

Split the data into *three* sets:

- **Training set:** the data on which the model is trained
- **Test set:** the data reserved for evaluation. New to the model
- **(Cross) Validation set:** reserve another set of data just for cross-validating across different models

e.g., 60/20/20 split: 60% training, 20% validation, 20% test



Training set ➡ fit the model parameters

Validation set ➡ choose among different models

Test set ➡ evaluate the performance of the final model

Computing training, validation, and test errors: linear regression example

Fit the model parameters by minimizing the cost function:

$$J = \underbrace{\frac{1}{m_{train}} \sum_{i=1}^{m_{train}} (y - f(x))^2}_{\text{MSE}_{train}} + \lambda \sum_{j=1}^n w_j^2$$

Training error:

$$J_{train} = \text{MSE}_{train}$$

(Cross) validation error (CV) for model selection:

$$J_{cv} = \text{MSE}_{cv}$$

Test error for model evaluation:

$$J_{test} = \text{MSE}_{test}$$

Model selection: Choosing model based on validation error

d : degree of polynomial

$$d = 1, f(x) = w_1x + b$$

$$d = 2, f(x) = w_1x + w_2x^2 + b$$

$$d = 3, f(x) = w_1x + w_2x^2 + w_3x^3 + b$$

...

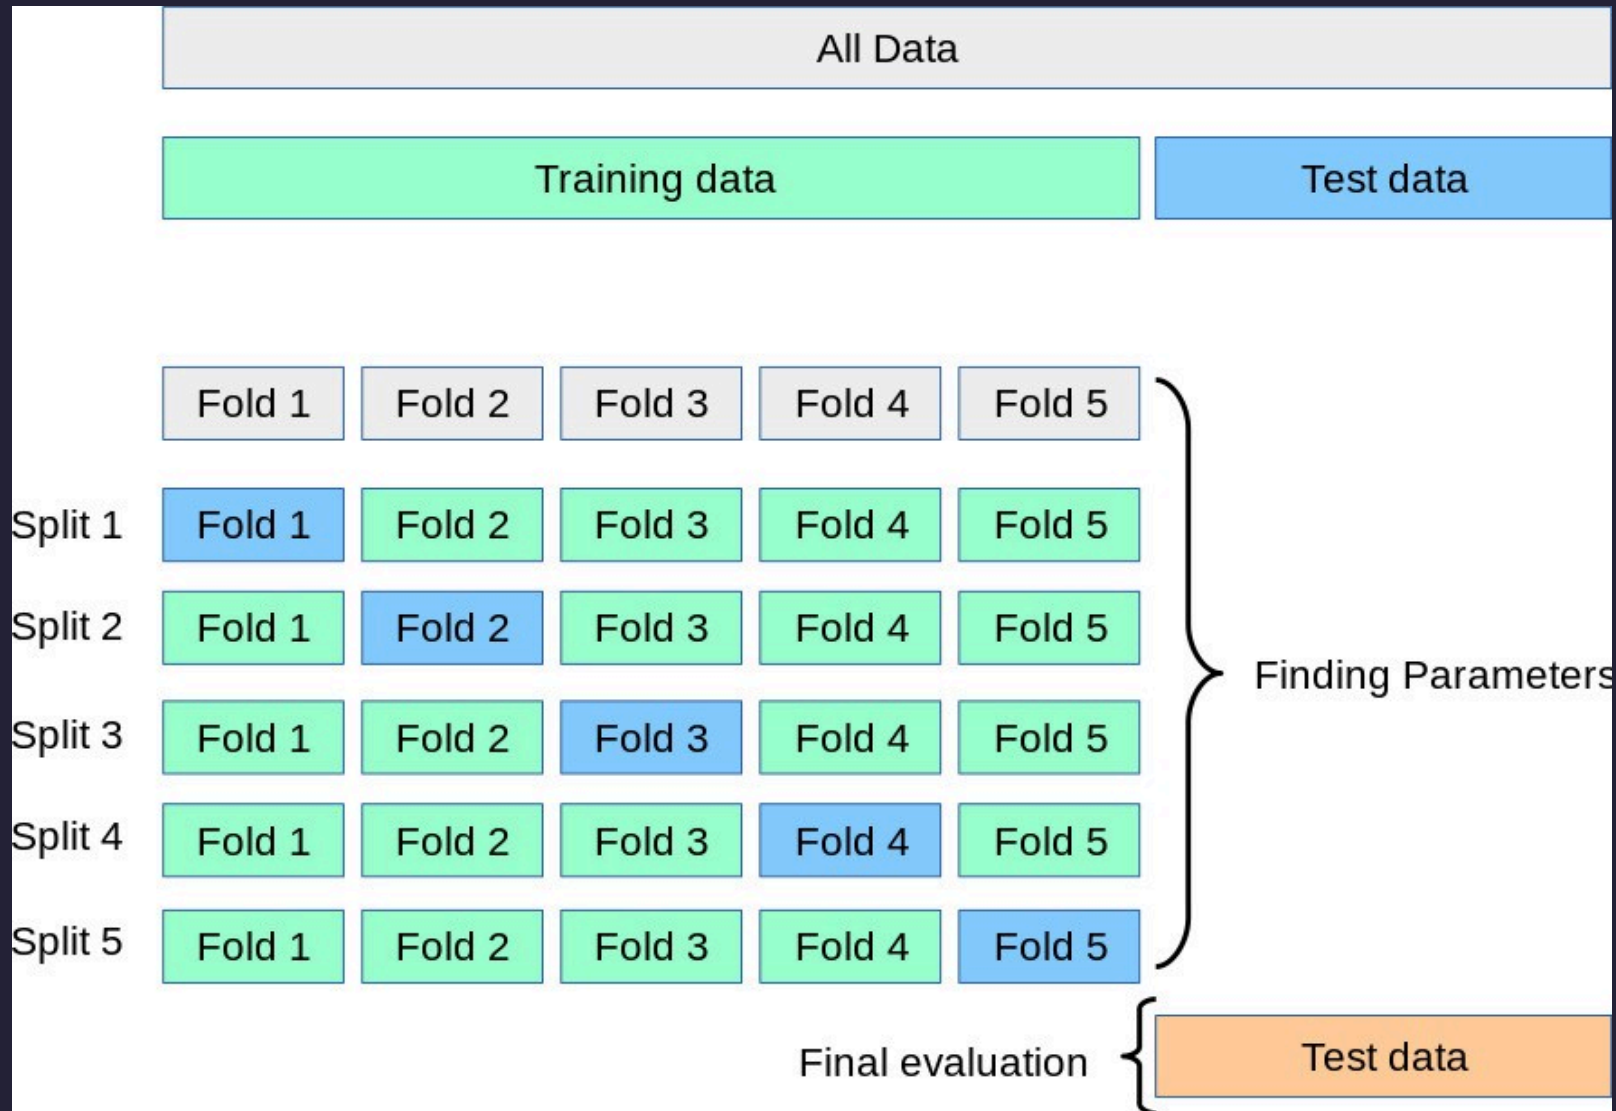
$$d = 10, f(x) = w_1x + w_2x^2 + \dots + w_{10}x^{10} + b$$

Calculate J_{cv} for each model (i.e., each d)

Choose the model with the lowest J_{cv}

Evaluate the final model using J_{test}

More efficient approach: k-fold cross-validation



k-fold cross validation

1. Split the data into k folds
2. Train the model on $k-1$ folds and validate on the remaining fold (compute validation error). Repeat k times, each time choosing a different fold as the validation set
3. Compute the average validation error (J_{cv})
4. Choose the configuration with the lowest J_{cv}
5. Retrain the model with the best configuration on the entire training dataset (m_{train})
6. Evaluate the final model using the test set (J_{test})

Hyperparameter tuning: choosing the optimal configuration based on J_{cv}

degree of polynomial ($d = 1, 2, 3$)

regularization ($\lambda = 0.01, 0.1, 1$)

learning rate

initial weights

number of layers in a neural network

number of hidden units in a neural network

...

Grid Search for hyperparameter tuning

Search a grid of hyperparameters

- all possible combinations

Lowest J_{cv} ➡ best hyperparameters

Pros: Systematic search

Cons: Computationally expensive


- More efficient search algorithms:
Random search, Bayesian
optimization, etc.

Grid Search

Hyperparameter_One = [a, b, c]


Hyperparameter_Two = [x, y, z]

Hyperparameter_X = [i, j, k]



Bias and variance

Bias and variance




High bias (underfit)

J_{train} is high (red)

J_{cv} is high (green)

$$J_{cv} \approx J_{train}$$



High variance (overfit)

J_{train} is low (red)

J_{cv} is high (green)


$$J_{cv} > J_{train}$$

Bias-variance tradeoff

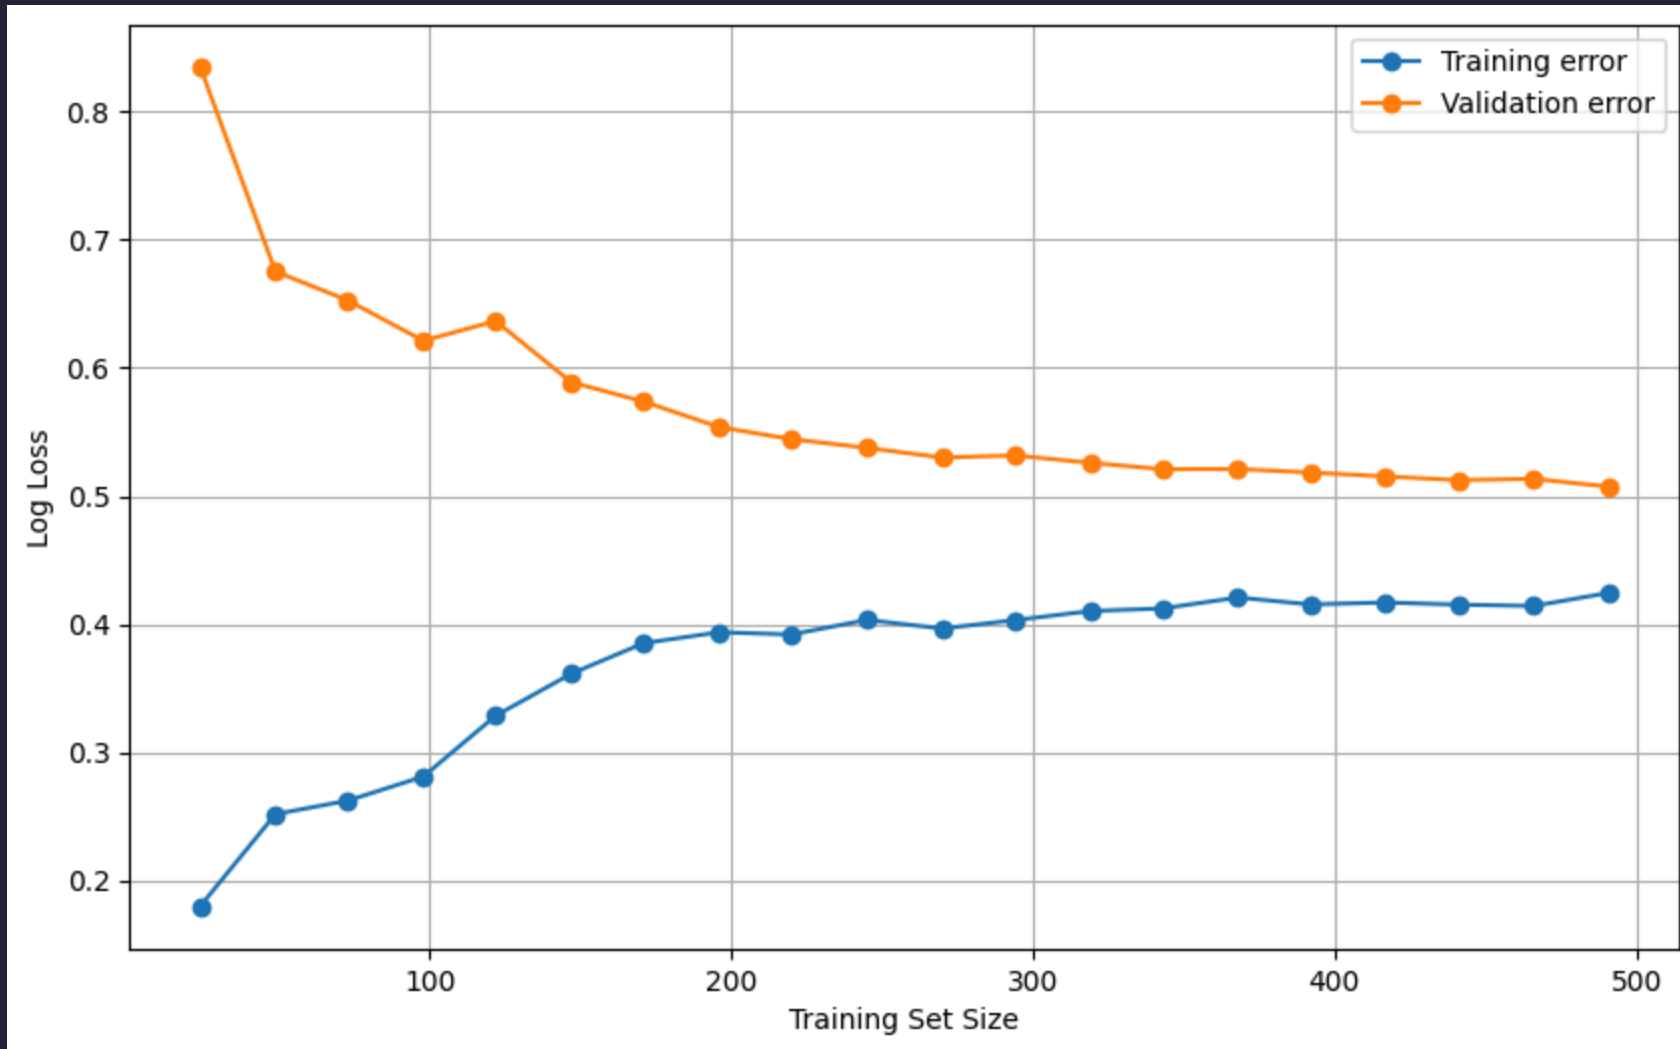
The balance between underfit and overfit

Optimal model complexity

- polynomial degree
- regularization parameter
- number of layers in a neural network
- ...



Learning curves: error vs. training set size




Will collecting more data help?

$$J_{train} < J_{cv}$$

➡ high variance (overfit)

More data helps



Will collecting more data help?


$$J_{train} \approx J_{cv}$$

$$J_{train} > J_{reference}$$

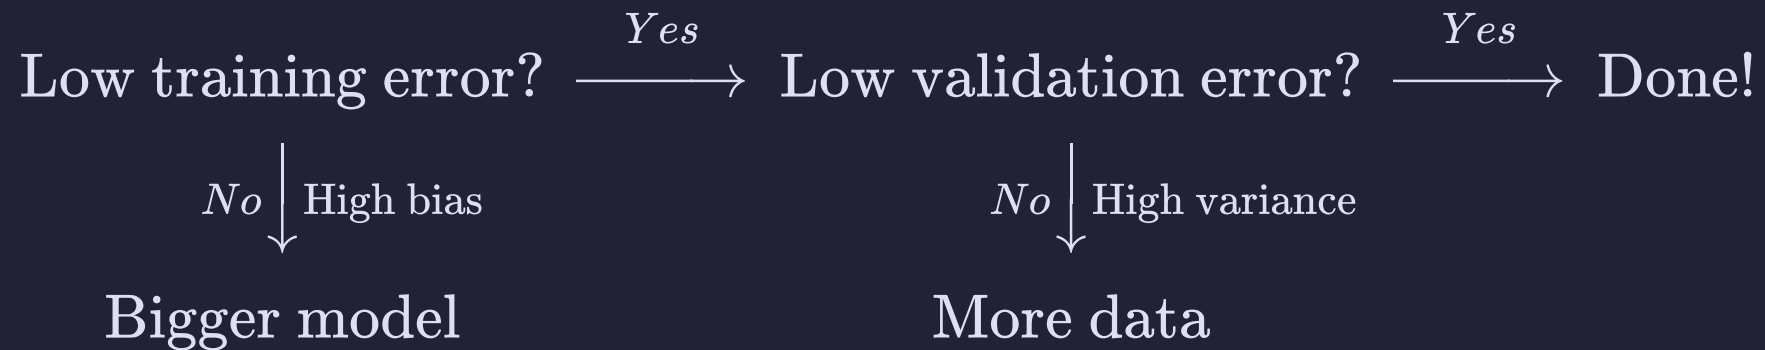
➡ high bias (underfit)

More data won't help

Fit a more complex model



Debugging your ML model: High bias or high variance?



Does it fix bias or variance problem?

$$J = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)}) + \lambda \sum_{j=1}^n w_j^2$$

Got large prediction errors. What do you do?

- Get more training examples
- Try smaller sets of features
- Try getting additional features
- Try adding polynomial features (e.g., x_1^2 , $x_1 x_2$, x_2^2)
- Try decreasing λ
- Try increasing λ

Grid search for hyperparameter tuning